CURRICULUM VITAE

Vladimiras Levinas

mob. **+37061428752**

E-mail **levin@linux.com**

Vilnius, Lithuania

OBJECTIVE

Lead DevOps Engineer with 15+ years of hands-on experience designing, building, and operating highly available, secure, and scalable production platforms in regulated enterprise, fintech, banking, and government environments. Proven track record in leading cloud-native and hybrid data platforms handling large-scale analytical and transactional workloads. Strong expertise in Linux systems, Kubernetes, containerization, GitOps, CI/CD automation, Infrastructure as Code, observability, and SRE practices. Seeking a leadership role with platform ownership responsibility, focused on operational excellence, reliability, and continuous improvement of mission-critical systems. List of Projects

CORE SKILLS

Operating Systems & Networking
Linux (Ubuntu, Debian, CentOS), TCP/IP, DNS, TLS, load balancing, firewalling

Containers & Orchestration
Docker, Docker Swarm, Kubernetes (k3s), Helm, KEDA, HPA

CI/CD & Automation
GitLab CI, Jenkins, GitHub Actions, Argo CD, GitOps workflows

Infrastructure as Code
Terraform, Ansible

Monitoring & Observability
Prometheus, Grafana, Alertmanager, OpenTelemetry, ELK Stack, Zabbix, Datadog

Cloud & Platforms
AWS, Hetzner, hybrid and on-prem environments

Databases & Data Platforms
SAP IQ, Oracle, PostgreSQL, SQL-based systems, ETL pipelines

Scripting & Languages
Bash, Python, SQL

AI/ML Infrastructure & Platforms

Self-hosted LLM infrastructure (Ollama), Model Context Protocol (MCP) servers,

Kubeflow (ML pipelines, distributed training), PyTorch (model deployment),

JupyterHub (multi-user notebooks)

Practices & Methodologies
SRE, ITIL, incident management, postmortems, capacity planning, ISO 9001, CMMI

Current Professional Certifications

- Linux Foundation Certified IT Associate (LFCA)

- Certified Kubernetes Administrator (CKA)

- Prometheus Certified Associate (PCA)

- OpenTelemetry Certified Associate (OTCA)

- HashiCorp Certified: Terraform Associate (003)

- ITIL Foundation Certificate in IT Service Management

- PMI Project Management Professional – PMBOK Framework

PROFESSIONAL EXPERIENCE

DBITS LT — Senior DevOps Engineer

Vilnius / International · 08/2007 – Present

Designed and modernized **data warehouse and ETL systems** by integrating **SAP**, **DevOps**, and **cloud-native** architectures for financial, retail, and government clients.

- Architected and deployed hybrid data platforms, integrating SAP-based ETL pipelines with microservice orchestration layers on Kubernetes.

- Containerized ETL components and services using Docker and orchestrated them through Kubernetes, achieving scalability and resilience.

- Automated infrastructure provisioning and configuration with Terraform and Ansible, ensuring reproducibility across Dev/Test/Prod environments.

- Developed CI/CD pipelines (GitLab, Jenkins) for automated deployment of ETL jobs, SQL updates, and monitoring agents.

- Implemented intelligent scaling mechanisms using HPA and KEDA for performance optimization under variable workloads.

- Integrated observability stack — Prometheus, Grafana, ELK, OpenTelemetry — enabling transparency and proactive issue detection.

- Automated backup and recovery for SAP IQ with encryption, integrity validation, and point-in-time recovery.

- Built self-hosted AI infrastructure platform on Kubernetes featuring Ollama for LLM inference, Model Context Protocol (MCP) servers for tool integration, and GitLab CI automation for model deployment workflows.

- Evaluated and tested ML platform technologies including Kubeflow for pipeline orchestration, PyTorch for model deployment, and JupyterHub for multi-user notebook environments.

C.T.Co 09/2005 – 05/2007 (Riga, Latvia)

Senior Business Analyst, PMO

Supported enterprise software projects for financial and reinsurance clients, connecting business and development teams.
Led **CMMI Level 4** and **ISO 9001** process alignment, business analysis, and internal workflow optimization.

Key Achievements:

- Implemented **Corporate Knowledge Management** system — developed training programs and documentation standards for PMs and analysts.

- Managed **Company Information & Control System** — integrated **JIRA ↔ MS Project**, HR/Skills DB, asset registry, and KM portal.

Tools: Atlassian JIRA, Microsoft Project, Confluence, SQL, HR/Skills DB, Knowledge Management System.

BHRS – Senior Project Manager / Local Representative

Tallinn, Estonia — 2003–2005

Represented **Micros-Fidelio (USA)** — a global provider of hospitality software and hardware solutions.
Managed **end-to-end system implementation** and infrastructure projects for major hotel chains in Lithuania (Le Meridien, Holiday Inn, Radisson SAS, Reval Lietuva) and provided support across the Baltics and Moscow.

Key Achievements:

- Led deployment and integration of **Micros-Fidelio PMS/POS systems** and back-office applications.

- Administered **Linux/Windows servers**, **Oracle/SQL databases**, and optimized network performance for high-traffic hospitality environments.

- Adapted systems to local business and legal requirements, ensuring compliance and stability.

- Coordinated cross-company implementations with partners including **Siemens** and **Comliet**.

Technologies: Linux/Windows Server, Oracle DB, Micros-Fidelio PMS/POS, TCP/IP networking, virtualization, backup and security tools.

Interlogics – Senior Programmer / Project Manager

Vilnius, Lithuania — 2000–2003

Developed and maintained **Net Symphony**, a proprietary web application server and scripting platform for enterprise **e-commerce** and **CRM** systems.
Led the **Vilnius development office** and coordinated projects for Japanese corporate clients such as **Sony**, **Mitsubishi**, and **Kajima**.

Key Achievements:

- Designed and implemented **frontend architecture** and administrative web interfaces.

- Built custom **JavaScript/XML extensions** for Dreamweaver and integrated **Net Symphony Script** into enterprise platforms.

- Delivered **personalization**, **CRM**, and **e-commerce modules**, improving scalability and usability.

- Developed **server-side components** in Java and optimized **Apache/Oracle** performance on Solaris/Linux.

Technologies: Java, JavaScript, XML, HTML/DHTML, Apache, Oracle DB, Sun Solaris, Linux, Net Symphony, Dreamweaver API.

Selected Projects

Banking Sector

1. Hybrid Data Platform for Financial Analytics

Role: Lead DevOps Engineer
Client: Kapital Bank
Period: 2019 – 2025

Led the design and implementation of a hybrid analytical data platform for a large commercial bank, combining **SAP IQ** with cloud-native infrastructure. Owned platform architecture with a focus on high availability, fault tolerance, and strict SLA compliance. Led the containerization of ETL workloads and their orchestration using Kubernetes to support dynamic scaling under variable transactional load. Established standardized Dev, Test, and Production environments with consistent configuration management. Designed and implemented automated deployment and upgrade pipelines to reduce operational risk. Acted as the primary technical owner responsible for platform stability and operational excellence.

Stack: Kubernetes (k3s), Docker, SAP IQ, Helm, Terraform, Ansible

2. Containerization of Legacy ETL Workloads

Role: Senior DevOps Engineer (Technical Lead)
Client: Respublika Bank
Period: 2017 – 2021

Led the migration of legacy, monolithic ETL systems to a modern container-based platform. Assessed existing SQL and batch workloads and redesigned them for secure and reliable execution in Docker containers. Defined Kubernetes resource limits, quotas, and operational guardrails to prevent performance degradation. Implemented CI/CD pipelines to automate ETL deployment and lifecycle management. Ensured backward compatibility with existing SAP-based data processing systems. Significantly improved system predictability, scalability, and operational maintainability.

Stack: Docker, Kubernetes, Bash, SQL, GitLab CI

---

3. GitOps-Based CI/CD Platform

Role: Lead DevOps Engineer
Client: Unibank
Period: 2020 – 2024

Designed and implemented a GitOps-based CI/CD platform for regulated banking environments. Established Git as the single source of truth for application and infrastructure configuration. Implemented Argo CD to automate cluster synchronization and enforce deployment consistency. Introduced mandatory code reviews, controlled promotion between environments, and standardized release processes. Enabled fast rollback and full auditability of deployments. Reduced manual operational errors and improved overall platform reliability.

Stack: Argo CD, GitLab CI, Helm, Kubernetes, Git

Government Big Data & Analytics Track

Role: Lead DevOps Engineer / Data Platform Architect
Clients:

- Ministry of Healthcare of the Republic of Kazakhstan

- Kazakhtelecom

- Ministry of Emergency Situations of the Republic of Kazakhstan

Period: 2013 – 2017

Led DevOps and platform engineering efforts for **national-scale government Big Data systems** across healthcare, telecommunications, and emergency response. Designed and operated centralized data platforms capable of processing **billions of transaction-level records**. Owned infrastructure reliability, scalability, and performance under sustained and peak loads. Solved complex operational challenges related to large-scale ingestion, analytical workloads, and long-running queries. Implemented forecasting and analytics platforms supporting capacity planning, risk assessment, and operational decision-making. Established standardized ETL pipelines, monitoring, backup, and operational practices for mission-critical systems.

Key Leadership Focus:

- Platform ownership and reliability at national scale
- High-availability systems for regulated environments
- Performance optimization for billion-record analytics
- Operational standards, observability, and incident readiness

Stack: Linux, SQL-based Data Warehouses (SAP IQ and similar), Kubernetes, Docker, ETL pipelines, CI/CD, Prometheus, Grafana, monitoring and backup systems

---

4. Self-Hosted AI Infrastructure Platform

Role: Lead DevOps Engineer

Environment: Personal Homelab / R&D

Period: 2024 – Present

Designed and deployed production-grade self-hosted AI infrastructure on Kubernetes for LLM operations and ML workflows. Built comprehensive platform combining large language models, custom tooling, and ML pipeline automation.

Key Components:

- Ollama deployment on K3s for self-hosted LLM inference (Llama, Mistral, CodeLlama)
- Model Context Protocol (MCP) server integration with GitLab, Kubernetes APIs
- CI/CD automation for model deployment and version management
- Complete observability stack (Prometheus, Grafana, Loki) for AI workload monitoring
- Resource management and GPU scheduling for ML workloads

Technical Implementation:

- Kubernetes-native architecture with GitOps workflows (ArgoCD)
- Automated model deployment pipelines via GitLab CI
- MCP servers providing LLM access to infrastructure APIs and data sources
- Secure API gateway for model inference endpoints
- Comprehensive logging and metrics for AI operations

ML Platform Experience:

- Evaluated and tested Kubeflow for ML pipeline orchestration
- PyTorch model containerization and deployment patterns

- JupyterHub integration for data science workflows
- Multi-tenant resource isolation for ML workloads

Impact:
- Demonstrated feasibility of enterprise-grade self-hosted AI infrastructure
- Established best practices for LLM deployment on Kubernetes
- Created reusable automation patterns for AI/ML operations
- Reduced dependency on external AI services

Stack: Kubernetes (K3s), Ollama, MCP, GitLab CI, ArgoCD, Prometheus, Grafana, Loki, Docker, Terraform, Ansible, Python, Bash

Skills Developed: AI/ML infrastructure, LLM operations, model deployment automation, GPU resource management, ML observability